

---

# Font Generation Trends In Deep Learning

---

## Machine Learning Winter School 2020

2020. 12.17

송실대학교

최재영

## **Contents**

1. Deep Learning
2. Computer Vision Tasks
3. Convolutional Neural networks (CNNs)
4. Generative Modeling
5. Generative Adversarial Networks (GANs)
6. Image to Image Translation (I2I)

# Overview

## Artificial Intelligence

Any technique that enables computers to mimic human behavior



## Machine Learning

Ability to learn without explicitly being programmed



## Deep Learning

Extract patterns from data using neural networks



## Why Deep Learning

Hand engineered features are time consuming and not scalable in practice  
Can we learn the *underlying features* directly from data?



## Why Deep Learning




*Traditional software development involves manually programming explicit rules.*

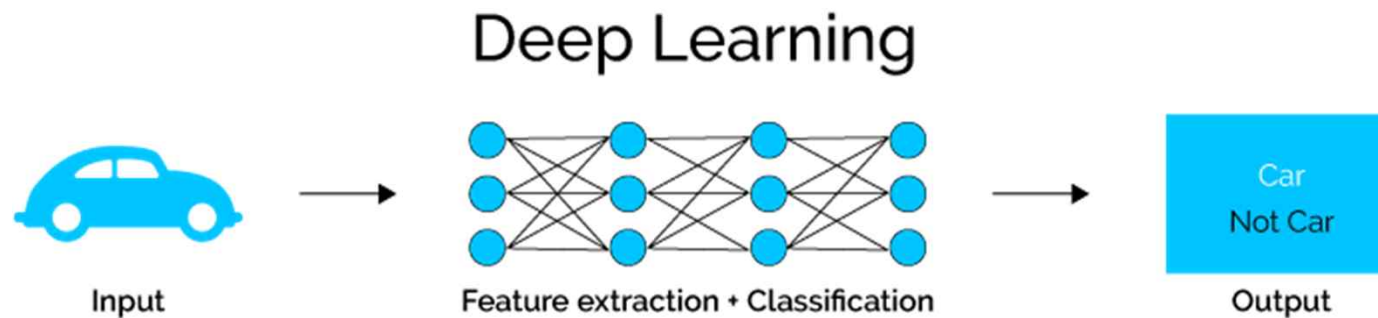


*Machine Learning is about **learning rules** from data.*

## Why Deep Learning



*Traditional machine learning uses hand-crafted features, which is tedious and costly to develop.*




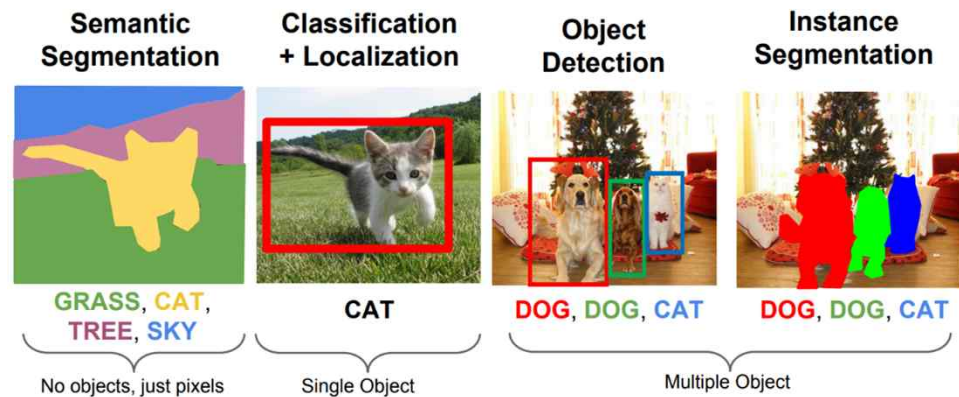
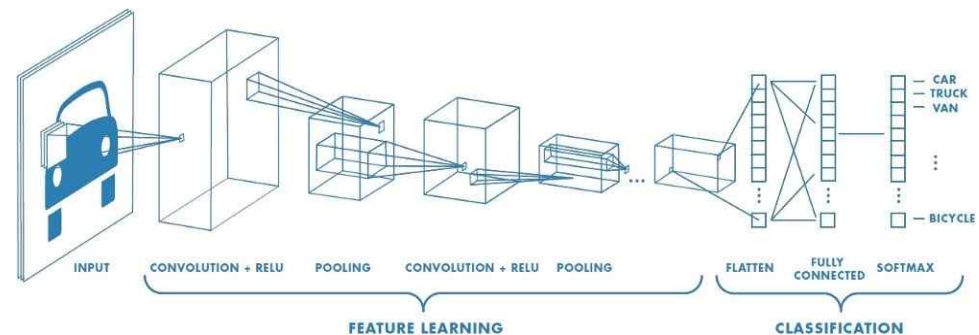
*Deep Learning learns **hierarchical** representations from the data itself, and scales with more data.*



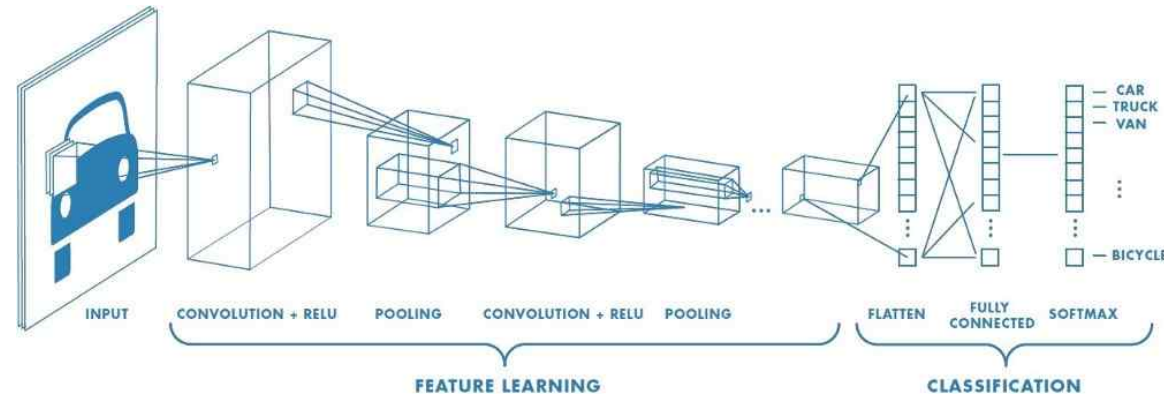
# Computer Vision Tasks

- Deep Learning has different architectures:
  - Multi-Layer Perceptron (MLP)
  - Convolutional neural net (CNN)
- Deep Learning introduces following computer vision application:
  - Image Classification
  - Semantic Segmentation
  - Object Classification + Localization
  - Object Detection
  - Instance Segmentation
  - Image Generation
- There are still many challenging problems to solve in computer vision.

“Single deep learning model can learn meaning from images and perform vision tasks, obviating the need for a pipeline of specialized and hand-crafted methods.”



# Computer Vision Tasks

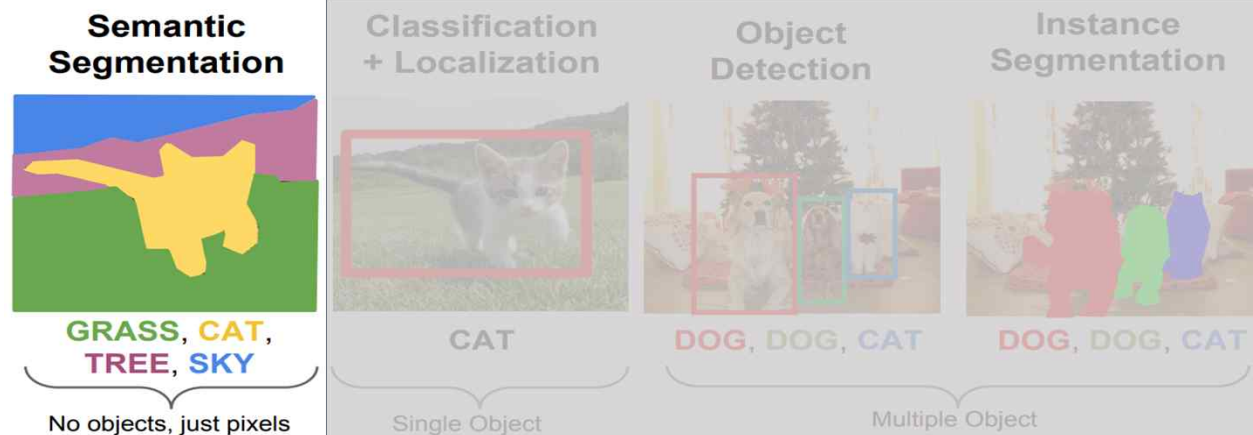


## Image Classification:

- The convolutional neural network, or **CNN** for short, is a **specialized** type of neural network model used for image classification tasks
- Image classification is a **supervised** learning problem
- It refers to the task of
  - **extracting** information from the input image and
  - **assigning** a label to an entire image or photograph from a fixed set of defined categories
- This problem is also referred to as
  - object classification and
  - image recognition

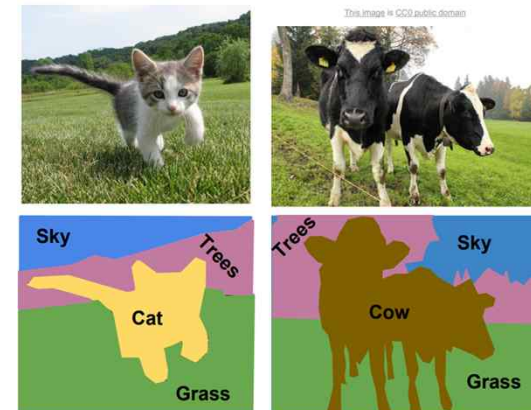


# Computer Vision Tasks

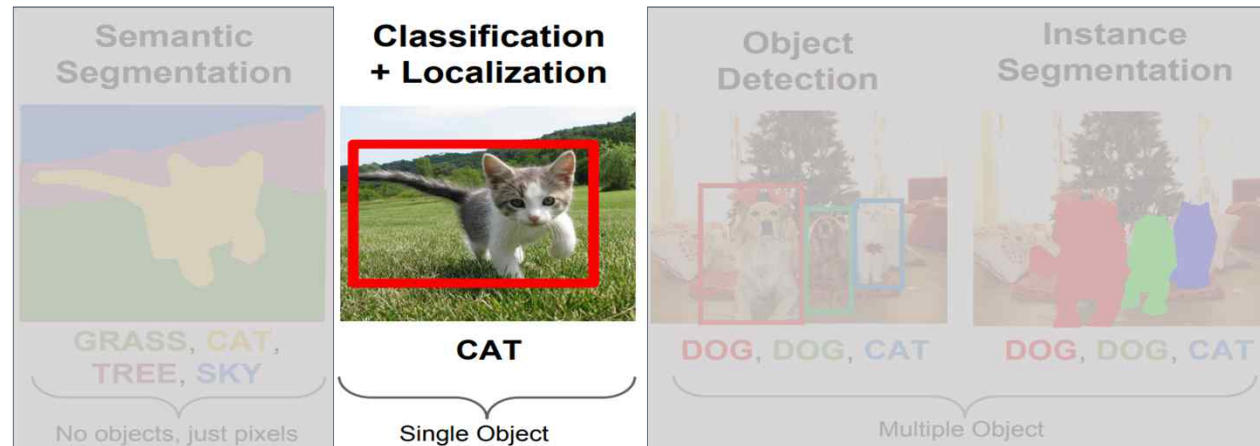


## Semantic Segmentation:

- There are two kinds of segmentation tasks in computer vision:
  - Semantic Segmentation
  - Instance Segmentation
- Label **each pixel** in the image with a **category label**
- Can't **differentiate** instance, only care about pixels
- This issue is fixed by **Instance Segmentation**



# Computer Vision Tasks



## Object Classification + Localization

- Image classification with localization involves:
  - **assigning** a class label to an image and
  - showing the **location of the object** in the image by a bounding box (drawing a box around the object)
- There should be only **single object** in the image, it cannot classify and localize **multiple objects** presented in the image
- This task is also called "**single-instance localization**"

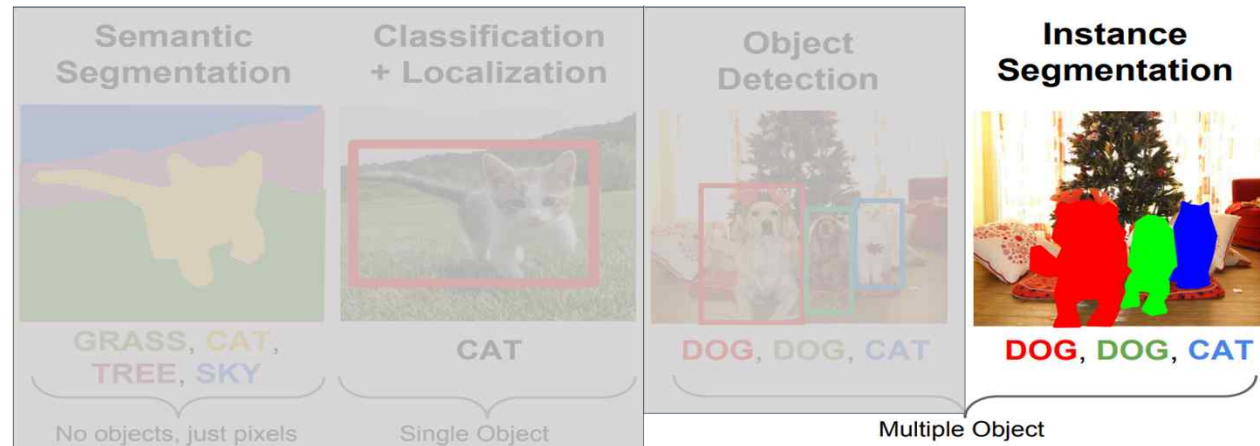
# Computer Vision Tasks



## Object Detection:

- Object Detection is the ability
  - to detect or identify **objects** in any given image
  - along with their **spatial position** in the given image
- It is commonly used in **applications** such as
  - image retrieval
  - security
  - surveillance, and
  - automated vehicle parking systems


# Computer Vision Tasks



## Instance Segmentation:

- Instance segmentation is an **important step** to achieving a comprehensive **image recognition** and object detection algorithms
- It has the ability to
  - identifies each instance of each object featured in the image
  - instead of categorizing each pixel like in **semantic segmentation**
- For example, instead of classifying five dogs as **one** instance, it will **identify** each **individual** dog in the image.

# Computer Vision Tasks



## Image Generation:

- Image generation is the task of
  - generating targeted modifications of existing images
  - or entirely new images set
- It involves automatically **discovering and learning the regularities or patterns** in input data
- This is a very broad area that is rapidly advancing by the use of **CNNs** and **Generative Adversarial Networks (GANs)**

## Generative Adversarial Nets – GAN

*Proposed an architecture that can learn the generative model of any data distribution through adversarial methods with excellent performance (2014).*

---

### Generative Adversarial Nets

---

**Ian J. Goodfellow, Jean Pouget-Abadie\*, Mehdi Mirza, Bing Xu, David Warde-Farley,  
Sherjil Ozair,<sup>†</sup> Aaron Courville, Yoshua Bengio<sup>‡</sup>**  
Département d'informatique et de recherche opérationnelle  
Université de Montréal  
Montréal, QC H3C 3J7

**Yann LeCun** described GANs as

"the most interesting idea in the last 10 years in Machine Learning."

(Cited 25,700 at Dec 2020)




# Generative Modeling

**Goal:** Take as input training samples from some distribution and learn a model that represents that distribution.



Input Samples

Training data  $\sim P_{\text{data}}(x)$



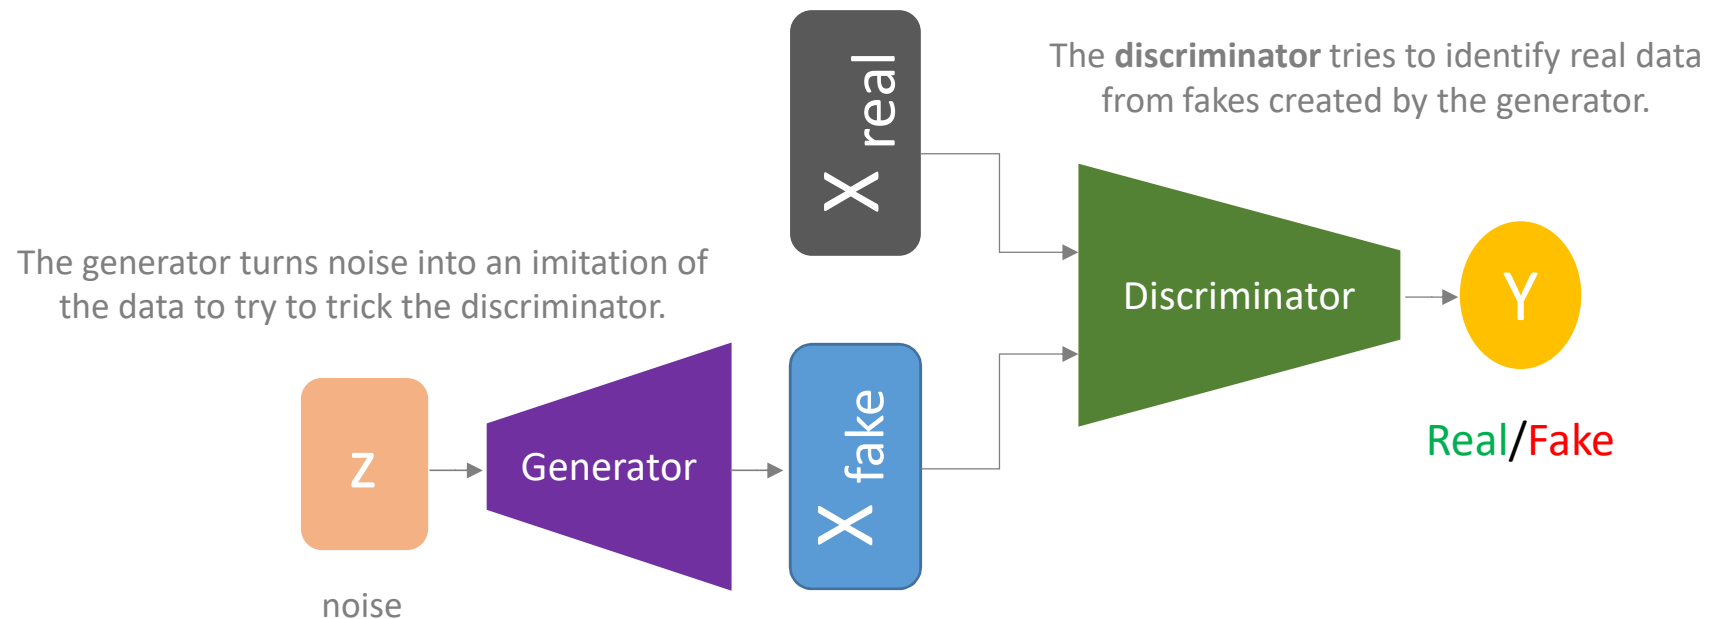
Generated Samples

Training data  $\sim P_{\text{model}}(x)$

How can we learn  $P_{\text{model}}(x)$  similar to  $P_{\text{data}}(x)$ ?

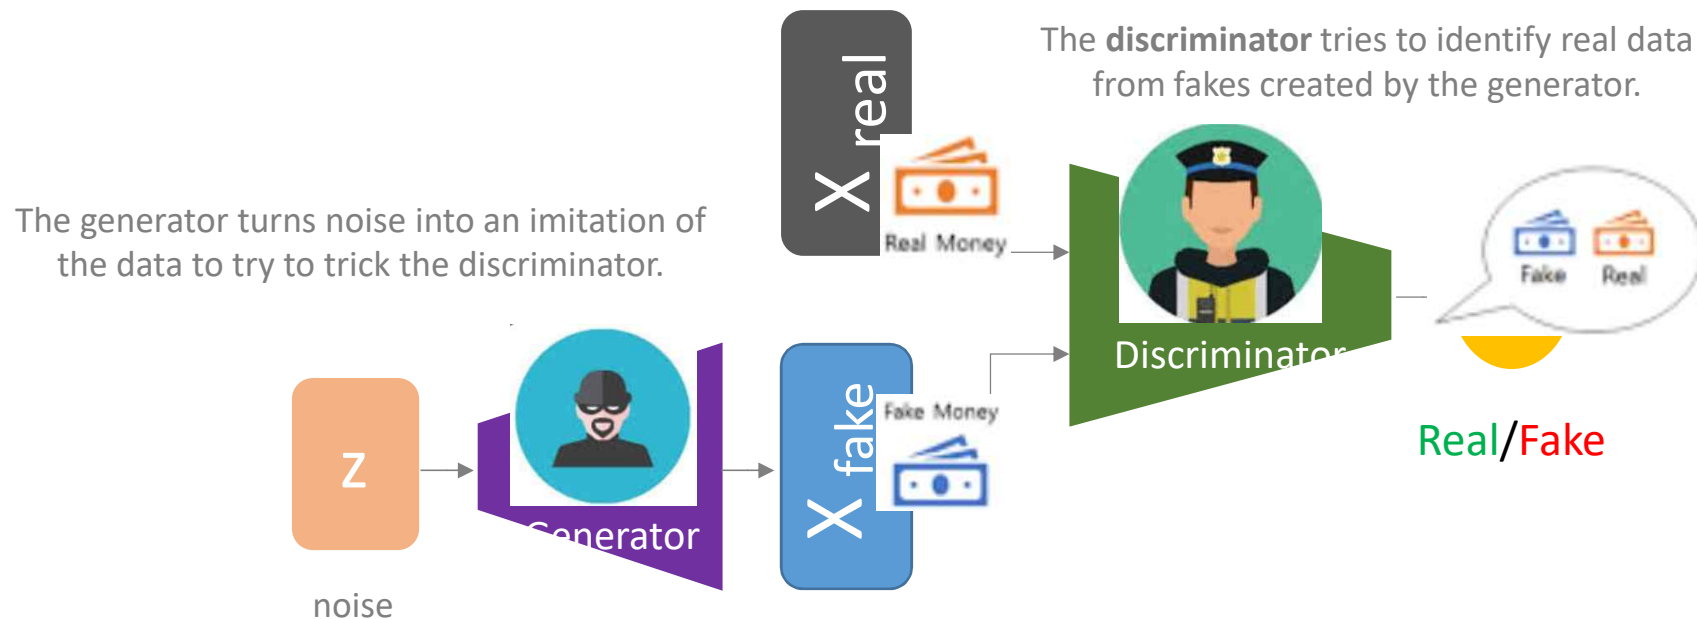
## Generative Adversarial Networks (GANs)

GANs are a way to make a generative model by having two neural networks compete with each other.




## Generative Adversarial Networks (GANs)

GANs are a way to make a generative model by having two neural networks compete with each other.



## Issues in Vanilla GAN

- Compared with other **generative models** such as Variational Autoencoders (VAEs),
  - images generated by GANs are usually **less blurred** and more **realistic**
- However, in practice, training GANs is **difficult** due to following reasons:
  - Do not have control on output
  - Training instability
- Various **extensions** of GANs have been proposed to **improve training stability**:
  - Conditional GANs (cGANs)
  - Deep Convolutional Generative Adversarial Networks (DCGANs)
  - Many More**



## Conditional Generative Adversarial Nets – cGAN

*Introduced the conditional version of generative adversarial nets.*

---

### Conditional Generative Adversarial Nets

---

**Mehdi Mirza**

Département d'informatique et de recherche opérationnelle

Université de Montréal

Montréal, QC H3C 3J7

`mirzamom@iro.umontreal.ca`

**Simon Osindero**

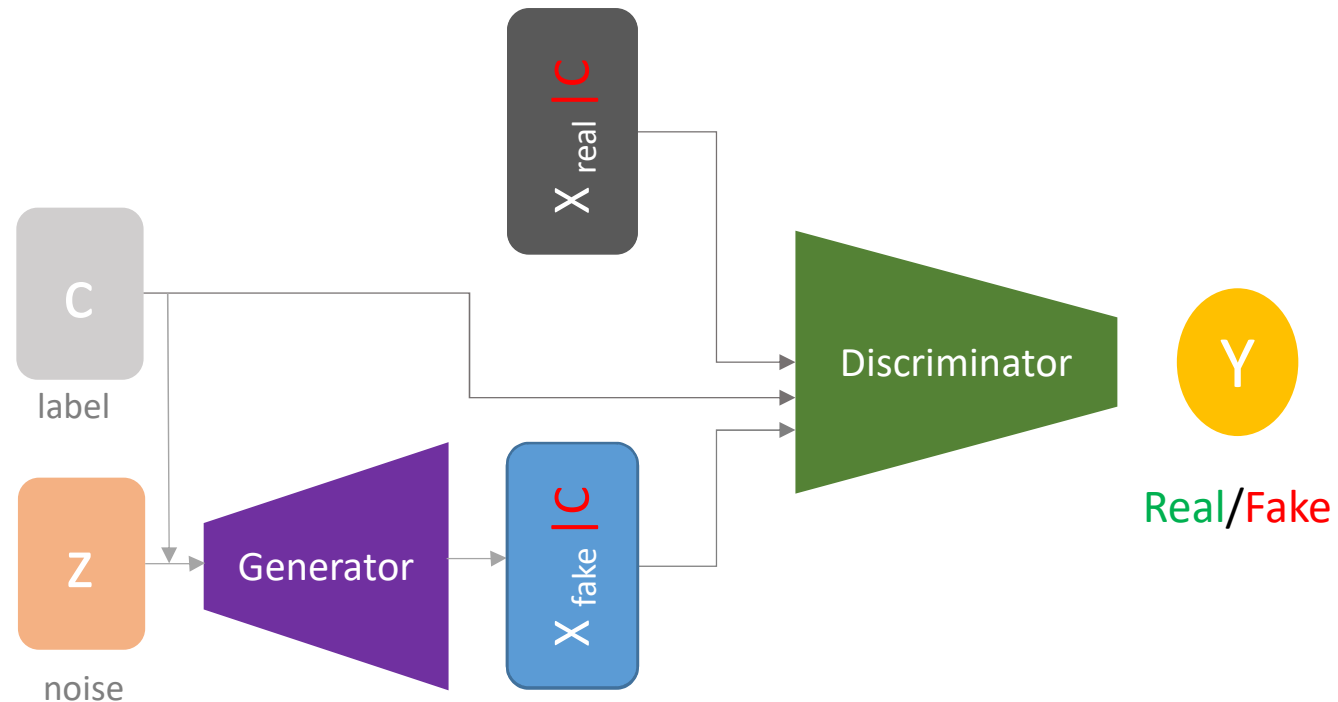
Flickr / Yahoo Inc.

San Francisco, CA 94103

`osindero@yahoo-inc.com`

## Conditional Generative Adversarial Networks (cGANs)

- In **cGANs**, the generator  $G$  and the discriminator  $D$  are conditioned on
  - some extra information **c**
- The extra information could be **class labels**, **text**, or **sketches**





## GANs VS cGANs Loss Functions

### Generative Adversarial Loss – GANs

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

*real samples*

*generated samples*

*Discriminator*

*correctly classify real  
data  $D(\mathbf{x}) \rightarrow 1$*

*correctly classify wrong  
data  $D(G(\mathbf{z})) \rightarrow 0$*

*Generator*

*Generate sample similar to  
real ones  $D(G(\mathbf{z})) \rightarrow 1$*

### Conditional Generative Adversarial Loss – cGANs

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x}|\mathbf{y})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z}|\mathbf{y})))]$$

*Same objective function as GANs only condition  $\mathbf{y}$  is added  
in both generator and discriminator to control the output*

## Noise-to-Image Translation – DCGAN

*Proposed an architecture that's more stable in training GAN's  
and more likely to converge. (in 2015) (**Without labels**)*

### UNSUPERVISED REPRESENTATION LEARNING WITH DEEP CONVOLUTIONAL GENERATIVE ADVERSARIAL NETWORKS

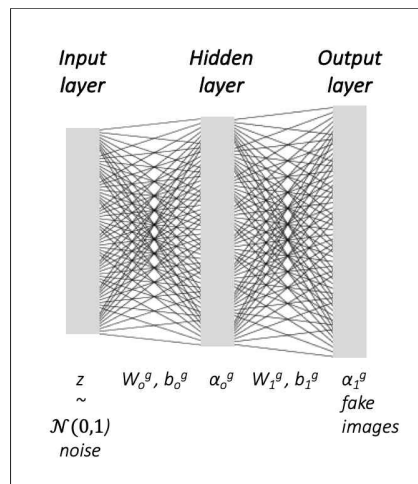
**Alec Radford & Luke Metz**  
indico Research  
Boston, MA  
`{alec,luke}@indico.io`

**Soumith Chintala**  
Facebook AI Research  
New York, NY  
`soumith@fb.com`


*After extensive model exploration we identified a family of architectures  
that resulted stable training across a range of datasets  
and allowed for training higher resolution and deeper generative models.*

# DCGAN Contribution

*DCGAN eliminates the need of fully connected layers in the network*



Generator




Discriminator


*DCGAN architecture uses a standard CNN architecture on the discriminative model*

## DCGAN Overall

Generator



Discriminator



# DCGAN Architectural Details

## Architectural Changes:

### 1. No Max Pooling:

- Replace max pooling layers with **strided convolutions** (discriminator) and **fractional-strided convolutions** (generator).

### 2. Using Batch Normalization

- Except the output layer for the **generator** and the input layer of the **discriminator**
- This mainly **tackles** two problems in **DCGAN** and in **deep neural networks** in general:
  - ✓ It normalizes the input to each unit of a layer.
  - ✓ It also helps to deal with poor initialization that may cause problems in gradient flow.

### 3. Remove Fully Connected Layers

- Remove fully connected hidden layers for deeper architecture

### 4. Use ReLU activation in Generator

- In all the layers of the **generator**, except for the **last one**.
- For the last convolutional layer, we will use **Tanh** activation function.


### 5. Use LeakyReLU activation in Discriminator

- Use LeakyReLU for all the convolutional layer after applying batch normalization.


$$\begin{array}{|c|c|c|c|c|c|} \hline 2 & 3 & 7 & 4 & 6 & 2 & 9 \\ \hline 6 & 6 & 9 & 8 & 7 & 4 & 3 \\ \hline 3 & 4 & 8 & 3 & 8 & 9 & 7 \\ \hline 7 & 8 & 3 & 6 & 6 & 3 & 4 \\ \hline 4 & 2 & 1 & 8 & 3 & 4 & 6 \\ \hline 3 & 2 & 4 & 1 & 9 & 8 & 3 \\ \hline 0 & 1 & 3 & 9 & 2 & 1 & 4 \\ \hline \end{array} \quad * \quad \begin{array}{|c|c|c|} \hline 3 & 4 & 4 \\ \hline 1 & 0 & 2 \\ \hline -1 & 0 & 3 \\ \hline \end{array} \quad = \quad \begin{array}{|c|c|c|} \hline 91 & 100 & 83 \\ \hline 69 & 91 & 127 \\ \hline 44 & 72 & 74 \\ \hline \end{array}$$

7\*7                      3\*3                      3\*3

$$\begin{array}{|c|c|c|c|c|c|} \hline 2 & 3 & 7 & 4 & 6 & 2 & 9 \\ \hline 6 & 6 & 9 & 8 & 7 & 4 & 3 \\ \hline 3 & 4 & 8 & 3 & 8 & 9 & 7 \\ \hline 7 & 8 & 3 & 6 & 6 & 3 & 4 \\ \hline 4 & 2 & 1 & 8 & 3 & 4 & 6 \\ \hline 3 & 2 & 4 & 1 & 9 & 8 & 3 \\ \hline 0 & 1 & 3 & 9 & 2 & 1 & 4 \\ \hline \end{array} \quad \begin{array}{|c|c|c|c|c|c|} \hline 2 & 3 & 7 & 4 & 6 & 2 & 9 \\ \hline 6 & 6 & 9 & 8 & 7 & 4 & 3 \\ \hline 3 & 4 & 8 & 3 & 8 & 9 & 7 \\ \hline 7 & 8 & 3 & 6 & 6 & 3 & 4 \\ \hline 4 & 2 & 1 & 8 & 3 & 4 & 6 \\ \hline 3 & 2 & 4 & 1 & 9 & 8 & 3 \\ \hline 0 & 1 & 3 & 9 & 2 & 1 & 4 \\ \hline \end{array} \quad \begin{array}{|c|c|c|c|c|c|} \hline 2 & 3 & 7 & 4 & 6 & 2 & 9 \\ \hline 6 & 6 & 9 & 8 & 7 & 4 & 3 \\ \hline 3 & 4 & 8 & 3 & 8 & 9 & 7 \\ \hline 7 & 8 & 3 & 6 & 6 & 3 & 4 \\ \hline 4 & 2 & 1 & 8 & 3 & 4 & 6 \\ \hline 3 & 2 & 4 & 1 & 9 & 8 & 3 \\ \hline 0 & 1 & 3 & 9 & 2 & 1 & 4 \\ \hline \end{array}$$




## DCGAN Model Architecture



## DCGAN Model Architecture - Generator

*DCGAN generator workflow – Step 1*




- Provide the generator a 100-dimensional noise vector as the input.
- After that, we project and reshape the input.



## DCGAN Model Architecture - Generator


*DCGAN generator workflow – Step 2*



- After feeding input vector, projection, and reshape:
  - We have **four convolutional operations** to be performed

## DCGAN Model Architecture - Generator


*DCGAN generator workflow – Step 3*



- Each time we get an **increment** in height and width
- At the same time, the **channels** keep on **reducing**.
- After the **first convolution operation**, we have **512** output channels.
- This keeps on reducing with each convolution operation from CONV1 to CONV4.
- After the third one, the **output** channels are **128**.

## DCGAN Model Architecture - Generator


*DCGAN generator workflow – Step 4*



- By the end, we have a generated image of 64x64 dimensions and three output channels.
- Except for the first convolution layer, all the other layers have a stride of 2.

## DCGAN Model Architecture - Discriminator


*DCGAN discriminator workflow*



- The discriminator takes an **image** as **input** from **generator**,
- Passes through **convolution** stacks, and
- Output a **probability** (sigmoid value) telling whether or not the image is **real**.

## Image-to-Image (I2I) Translation

I2I is a class of vision and graphics problems where the goal is to learn the mapping between an input image and an output image.




Input Image

Translated Image

## 121 Example use cases



Low-res to high-res




Blurry to sharp



Image to painting



LDR to HDR



Synthetic to real



Thermal to color



Day to night



Summer to winter



Noisy to clean



# Image-to-Image Translation – Pix2pix

- Class of computer vision and graphics problems
- Goal: Learn the **mapping** between input image and output image
  - OR learning the **translation** from input domain DA to output domain DB
- This can be achieved by a Conditional GAN (cGAN) **pix2pix**

## Image-to-Image Translation with Conditional Adversarial Networks

Phillip Isola

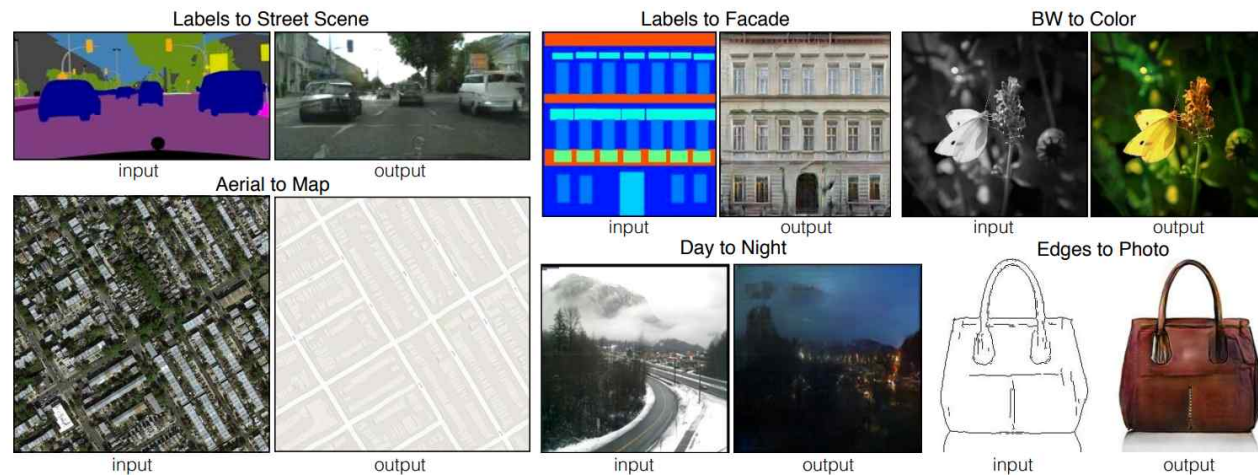
Jun-Yan Zhu

Tinghui Zhou

Alexei A. Efros

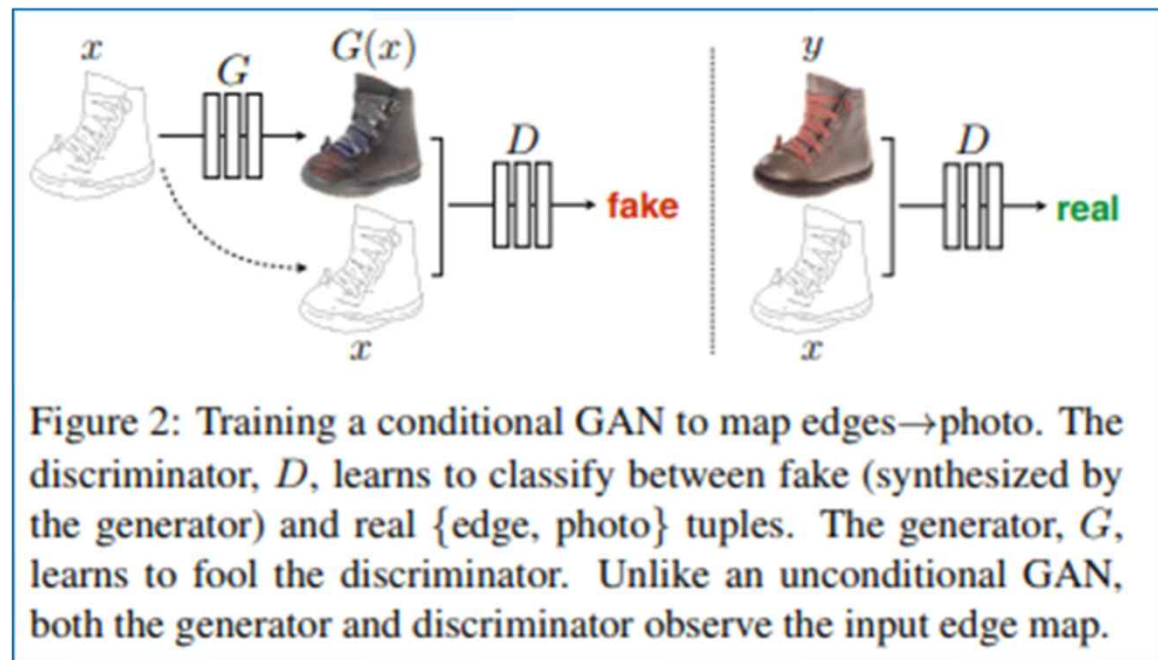
Berkeley AI Research (BAIR) Laboratory, UC Berkeley

{isola, junyanz, tinghuiz, efros}@eecs.berkeley.edu




## Pix2pix Architecture

- pix2pix depends on **paired** training data



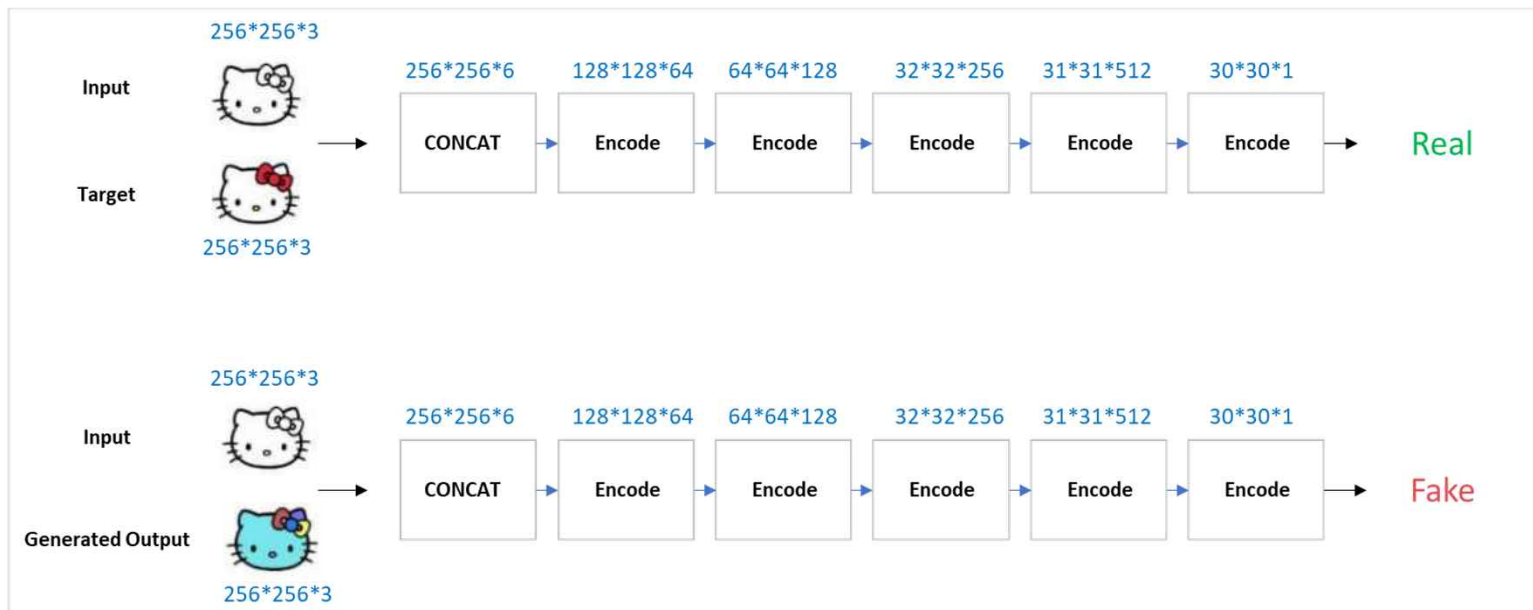
## Pix2pix Architecture – Generator

- The structure of the **generator** is called **U-Net** based encoder-decoder with **skip-connections**
- Job of generator is:
  - taking an input image (could be black and white image)
  - performing the transform to produce the target image (be a colored version)



## Pix2pix Architecture – Discriminator

- The structure of the **discriminator** is based on PatchGAN
- The Discriminator has the job of taking **two images**:
  - an input image and
  - an unknown image (which will be either a target or output image from the generator) and
  - decide if the other image was produced by the generator or not



## Pix2pix Objective Function

- Composite Adversarial and L1 Loss

*Conditional Generative Adversarial Loss – cGANs*

$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x,y} [\log D(x, y)] + \mathbb{E}_{x,z} [\log(1 - D(x, G(x, z)))],$$

*L1 loss*

$$\mathcal{L}_{L1}(G) = \mathbb{E}_{x,y,z} [\|y - G(x, z)\|_1].$$

*Combined Loss = cGANs + Lambda \* L1 loss*

$$G^* = \arg \min_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G)$$

## GANs: Recent Advances

Two **imaginary celebrities** that were dreamed up by a random number generator.

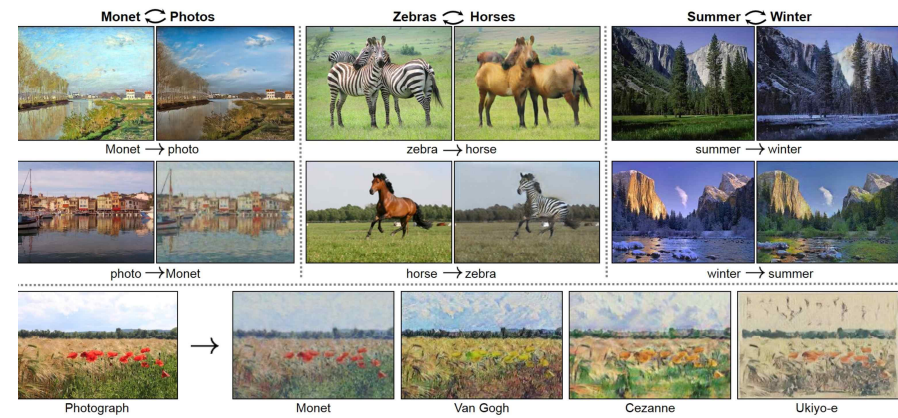
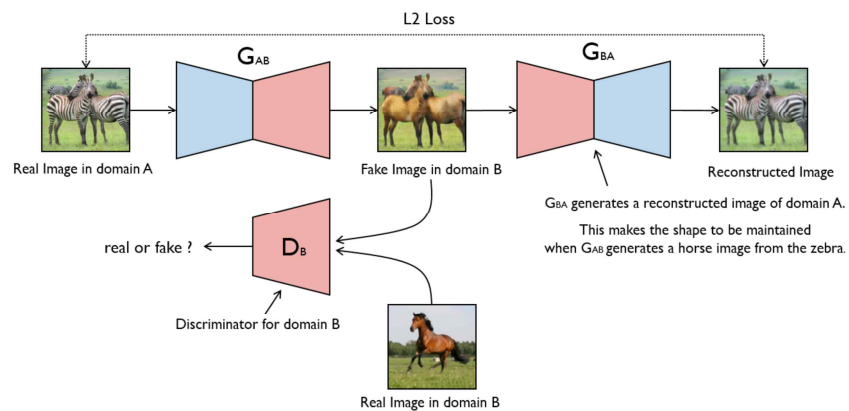


**“Progressive Growing of GANs for Improved Quality, Stability, and Variation”** by  
Tero Karras (NVIDIA), Timo Aila (NVIDIA), Samuli Laine (NVIDIA), Jaakko Lehtinen (NVIDIA and Aalto University)



# GANs: Recent Advances

## Domain transformation (CycleGAN)



**“Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks”** by

Jun-Yan Zhu\*, Taesung Park, Phillip Isola, Alexei A. Efros

**“Image-to-Image Translation with Conditional Adversarial Nets”** by

Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros

## GANs: Recent Advances

### Text-to-Image Translation (Text2image)

“The petals of the flower are pink in color and have a yellow center”



**“Generative Adversarial Text to Image Synthesis”** by

Scott Reed, Zeynep Akata, Xincheng Yan, Lajanugen Logeswaran



## Hands-on Lab

**Speaker:** Ammar Ul Hassan

### ✓ Font Generation using GANs

- Lab0- Setting up environment for running labs
- Lab1 - Building vanilla GAN in TensorFlow for MNIST
- Lab2 - Building DCGAN in TensorFlow for font dataset
- Lab3 - Building CDCGAN in TensorFlow for font dataset
- Lab4 - Building pix2pix in TensorFlow for font dataset

### ✓ Our current research on font generation

**Thank you**  

---

**감사합니다**